




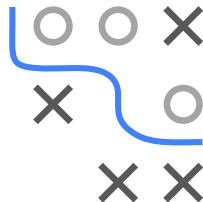
## Momentum on quadratic functions



- Effect of  $\varphi$



# RECAP: DYNAMICS OF MOMENTUM

- Let's investigate the role of  $\varphi$ .
- We can think of gradient descent with momentum as a damped harmonic oscillator: a weight on a spring. We pull the weight down and study the path back to the equilibrium in phase space (looking at the position and the velocity).
- Depending on the choice of  $\varphi$ , the rate of return to the equilibrium position is affected.



# RECAP: DYNAMICS OF MOMENTUM / 2


$$f : \mathbb{R} \rightarrow \mathbb{R}, x \mapsto -\frac{1}{2} \cos(x) \exp\left(\frac{1}{4}(x+2)^2\right) + 3, \quad x \in [-2, 2], \quad \alpha = 0,05$$




●  $\phi = 0.0$  ●  $\phi = 0.1$  ●  $\phi = 0.5$  ●  $\phi = 0.65$

## ©

[illegible]





Optimization in Machine Learning – 3 / 6




©

# RECAP: DYNAMICS OF MOMENTUM / 5

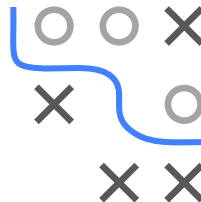

$$f : \mathbb{R} \rightarrow \mathbb{R}, x \mapsto \frac{1}{2}x^6 + \frac{3}{2}x^5 + 2x^3 + 5x^2 - 3x, \quad x \in \left[-\frac{9}{2}, 2\right], \quad \alpha = 0,02$$



●  $\phi = 0.0$  ●  $\phi = 0.1$  ●  $\phi = 0.5$  ●  $\phi = 0.65$



# RECAP: DYNAMICS OF MOMENTUM / 6



Left: If  $\varphi$  is too large, we are underdamping. The spring oscillates back and forth and misses the optimum.


Middle: The best value of  $\varphi$  lies in the middle.

Right: If  $\varphi$  is too small, we are overdamping, meaning that the spring experiences too much friction and stops before reaching the equilibrium.

# Optimization in Machine Learning

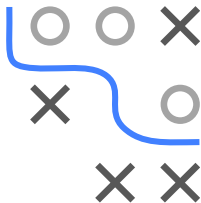
## First order methods

### SGD



#### Learning goals

- SGD
- Stochasticity
- Convergence
- Batch size






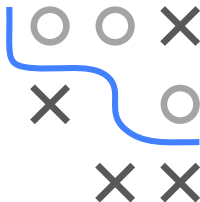
# STOCHASTIC GRADIENT DESCENT

**Issue:** Data-sets might be very large and gradients expensive to evaluate


**Idea:** Use a smaller (random) subset of data-points to evaluate gradient and objective-function





●  $n = 5$  ●  $n = 3$  ●  $n = 2$  ●  $n = 1$





# STOCHASTIC GRADIENT DESCENT / 2



●  $n = 5$  ●  $n = 3$  ●  $n = 2$  ●  $n = 1$





# STOCHASTIC GRADIENT DESCENT / 4

Analysis of curvature:

Least-Squares objective function:

$$f(\theta) = \min_{\theta \in \mathbb{R}^d} \sum_{i=1}^n \left( \theta^\top \mathbf{x}^{(i)} - y^{(i)} \right)^2 = (\mathbf{X}\theta - \mathbf{y})^\top (\mathbf{X}\theta - \mathbf{y})$$


1st derivative of objective function:

$$\nabla f(\theta) = \frac{1}{2} (\mathbf{X}^\top \mathbf{X}) \theta - \mathbf{X}^\top \mathbf{y}$$

2nd derivative of objective function:

$$\nabla^2 f(\theta) = \frac{1}{2} (\mathbf{X}^\top \mathbf{X})$$

⇒ Hessian is only dependent of  $x$ -coordinates of data-points



# STOCHASTIC GRADIENT DESCENT / 5

Hessian matrix:

$$\mathbf{H} = \nabla^2 f(\theta) = \frac{1}{2} (\mathbf{X}^T \mathbf{X}) = \frac{1}{2} \begin{bmatrix} 4.97 & 3.3 \\ 3.3 & 2.5 \end{bmatrix}$$

Eigenvalues:

$$\det(\mathbf{H} - \mathbf{I}\lambda) = 0$$


solve for  $\lambda$ :

$$(H_{11} - \lambda)(H_{22} - \lambda) - H_{21}H_{12} = \lambda^2 - 7.47\lambda + 1.535 = 0$$

$$\lambda_1 = 7.2585, \quad \lambda_2 = 0.2115$$

$$\kappa(\mathbf{H}) = \lambda_1 / \lambda_2 \approx 34.3191$$

$\Rightarrow$  Rather high condition-number compared to matrix-values, positive definite matrix



# STOCHASTIC GRADIENT DESCENT / 6


Clustered data-points increase condition-number for Least-Squares method:

$$x_{old}^{(i)} = (0.7, 0.8, 1.5, 1.6, 2.0)$$



$$x_{new}^{(i)} = (0.5, 0.8, 1.2, 1.6, 2.0)$$

$$\kappa(\mathbf{H}_{new}) = \lambda_1/\lambda_2 \approx 24.6074$$

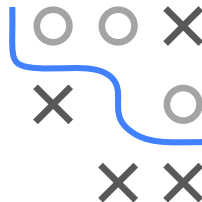
⇒ More evenly distributed data-points improve conditioning



# STOCHASTIC GRADIENT DESCENT / 7



●  $n = 5$  ●  $n = 3$  ●  $n = 2$  ● 1



# STOCHASTIC GRADIENT DESCENT / 8

NB: We use  $g$  instead of  $f$  as objective, bc.  $f$  is used as model in ML.

$g : \mathbb{R}^d \rightarrow \mathbb{R}$  objective,  $g$  **average over functions**:

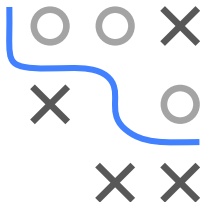
$$g(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n g_i(\mathbf{x}), \quad g \text{ and } g_i \text{ smooth}$$

Stochastic gradient descent (SGD) approximates the gradient

$$\begin{aligned}\nabla_{\mathbf{x}} g(\mathbf{x}) &= \frac{1}{n} \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) \quad := \quad \mathbf{d} \quad \text{by} \\ \frac{1}{|J|} \sum_{j \in J} \nabla_{\mathbf{x}} g_j(\mathbf{x}) &:= \hat{\mathbf{d}},\end{aligned}$$

with random subset  $J \subset \{1, 2, \dots, n\}$  of gradients called **mini-batch**.

This is done e.g. when computing the true gradient is **expensive**.

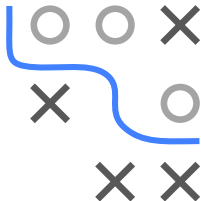




# STOCHASTIC GRADIENT DESCENT / 9

## Algorithm Basic SGD pseudo code

```
1: Initialize  $\mathbf{x}^{[0]}$ ,  $t = 0$ 
2: while stopping criterion not met do
3:   Randomly shuffle indices and partition into minibatches  $J_1, \dots, J_K$  of size  $m$ 
4:   for  $k \in \{1, \dots, K\}$  do
5:      $t \leftarrow t + 1$ 
6:     Compute gradient estimate with  $J_k$ :  $\hat{\mathbf{d}}^{[t]} \leftarrow \frac{1}{m} \sum_{i \in J_k} \nabla_{\mathbf{x}} g_i(\mathbf{x}^{[t-1]})$ 
7:     Apply update:  $\mathbf{x}^{[t]} \leftarrow \mathbf{x}^{[t-1]} - \alpha \cdot \hat{\mathbf{d}}^{[t]}$ 
8:   end for
9: end while
```



- Instead of drawing batches randomly we might want to go through the  $g_i$  sequentially (unless  $g_i$  are sorted in any way)
- Updates are computed faster, but also more stochastic:
  - In the simplest case, batch-size  $m := |J_k|$  is set to  $m = 1$
  - If  $n$  is a billion, computation of update is a billion times faster
  - **But** (later): Convergence rates suffer from stochasticity!



## SGD IN ML / 2

For large data sets, computing the exact gradient

$$\mathbf{d} = \frac{1}{n} \sum_{i=1}^n \nabla_{\theta} L \left( y^{(i)}, f \left( \mathbf{x}^{(i)} \mid \theta \right) \right)$$

may be expensive or even infeasible to compute and is approximated by


$$\hat{\mathbf{d}} = \frac{1}{m} \sum_{i \in J} \nabla_{\theta} L \left( y^{(i)}, f \left( \mathbf{x}^{(i)} \mid \theta \right) \right),$$

for  $J \subset 1, 2, \dots, n$  random subset.

**NB:** Often, maximum size of  $J$  technically limited by memory size.



# STOCHASTICITY OF SGD



Minimize  $g(x_1, x_2) = 1.25(x_1 + 6)^2 + (x_2 - 8)^2$ .

**Left:** GD. **Right:** SGD. Black line shows average value across multiple runs.

(Source: Shalev-Shwartz et al., Understanding Machine Learning, 2014.)

# STOCHASTICITY OF SGD / 2

Assume batch size  $m = 1$  (statements also apply for larger batches).

- **(Possibly) suboptimal direction:** Approximate gradient  $\hat{\mathbf{d}} = \nabla_{\mathbf{x}} g_i(\mathbf{x})$  might point in suboptimal (possibly not even a descent!) direction
- **Unbiased estimate:** If  $J$  drawn i.i.d., approximate gradient  $\hat{\mathbf{d}}$  is an unbiased estimate of gradient  $\mathbf{d} = \nabla_{\mathbf{x}} g(\mathbf{x}) = \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x})$ :

$$\begin{aligned}\mathbb{E}_i [\nabla_{\mathbf{x}} g_i(\mathbf{x})] &= \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) \cdot \mathbb{P}(i = i) \\ &= \sum_{i=1}^n \nabla_{\mathbf{x}} g_i(\mathbf{x}) \cdot \frac{1}{n} = \nabla_{\mathbf{x}} g(\mathbf{x}).\end{aligned}$$

**Conclusion:** SGD might perform single suboptimal moves, but moves in “right direction” **on average**.



# CONVERGENCE OF SGD

As a consequence, SGD has worse convergence properties than GD.

**But:** Can be controlled via **increasing batches** or **reducing step size**.


**The larger the batch size  $m$**

- the better the approximation to  $\nabla_{\mathbf{x}}g(\mathbf{x})$
- the lower the variance
- the lower the risk of performing steps in the wrong direction

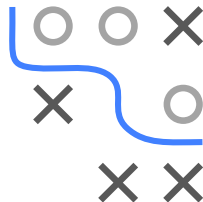

**The smaller the step size  $\alpha$**

- the smaller a step in a potentially wrong direction
- the lower the effect of high variance

As maximum batch size is usually limited by computational resources (memory), choosing the step size is crucial.



# EFFECT OF BATCH SIZE





SGD for a NN with batch size  $\in \{0.5\%, 10\%, 50\%\}$  of the training data.  
The higher the batch size, the lower the variance.

# Optimization in Machine Learning

## First order methods

## SGD Further Details




### Learning goals

- Decreasing step size for SGD
- Stopping rules
- SGD with momentum



# SGD WITH CONSTANT STEP SIZE

**Example:** SGD with constant step size.



Fast convergence of SGD initially. Erratic behavior later (variance too big).

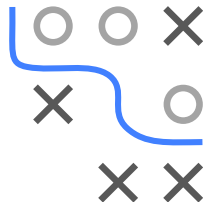


# SGD WITH DECREASING STEP SIZE

- **Idea:** Decrease step size to reduce magnitude of erratic steps.
- **Trade-off:**
  - if step size  $\alpha^{[t]}$  decreases slowly, large erratic steps
  - if step size decreases too fast, performance is impaired



# SGD WITH DECREASING STEP SIZE / 2

- Popular solution: step size fulfilling  $\alpha^{[t]} = \alpha^{[0]}/t$ .



Example continued. Step size  $\alpha^{[t]} = 0.2/t$ .

- Often not working well in practice: step size gets small quite fast.
- Alternative:  $\alpha^{[t]} = \alpha^{[0]}/\sqrt{t}$

# ADVANCED STEP SIZE CONTROL

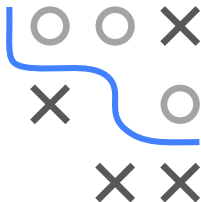
## Why not Armijo-based step size control?

- Backtracking line search or other approaches based on Armijo rule usually not suitable: Armijo condition

$$g(\mathbf{x} + \alpha \mathbf{d}) \leq g(\mathbf{x}) + \gamma_1 \alpha \nabla g(\mathbf{x})^\top \mathbf{d}$$

requires evaluating full gradient.

- But SGD is used to *avoid* expensive gradient computations.
- Research aims at finding inexact line search methods that provide better convergence behaviour, e.g., Vaswani et al., *Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates*. NeurIPS, 2019.





# STOPPING RULES FOR SGD

- **For GD:** We usually stop when gradient is close to 0 (i.e., we are close to a stationary point)
- **For SGD:** individual gradients do not necessarily go to zero, and we cannot access full gradient.
- Practicable solution for ML:
  - Measure the validation set error after  $T$  iterations
  - Stop if validation set error is not improving




# SGD AND ML

## General remarks:

- SGD is a variant of GD
- SGD particularly suitable for large-scale ML when evaluating gradient is too expensive / restricted by computational resources
- SGD and variants are the most commonly used methods in modern ML, for example:
  - Linear models

Note that even for the linear model and quadratic loss, where a closed form solution is available, SGD might be used if the size  $n$  of the dataset is too large and the design matrix does not fit into memory.
  - Neural networks
  - Support vector machines
  - ...



# SGD WITH MOMENTUM

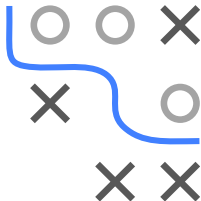
SGD is usually used with momentum due to reasons mentioned in previous chapters.

---

**Algorithm** Stochastic gradient descent with momentum

---

- 1: **require** step size  $\alpha$  and momentum  $\varphi$
  - 2: **require** initial parameter  $\mathbf{x}$  and initial velocity  $\boldsymbol{\nu}$
  - 3: **while** stopping criterion not met **do**
  - 4:     Sample mini-batch of  $m$  examples
  - 5:     Compute gradient estimate  $\nabla \hat{g}(\mathbf{x})$  using mini-batch
  - 6:     Compute velocity update:  $\boldsymbol{\nu} \leftarrow \varphi \boldsymbol{\nu} - \alpha \nabla \hat{g}(\mathbf{x})$
  - 7:     Apply update:  $\mathbf{x} \leftarrow \mathbf{x} + \boldsymbol{\nu}$
  - 8: **end while**
- 

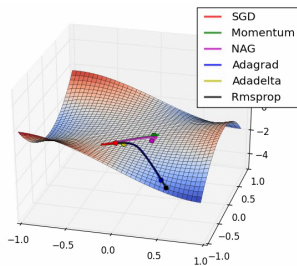




# Optimization in Machine Learning

## First order methods

## Adam and friends

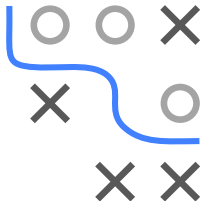



### Learning goals

- Adaptive step sizes
- Adam

## ADAPTIVE STEP SIZES

- Step size is probably the most important control parameter
- Has strong influence on performance
- Natural to use different step size for each input individually and automatically adapt them

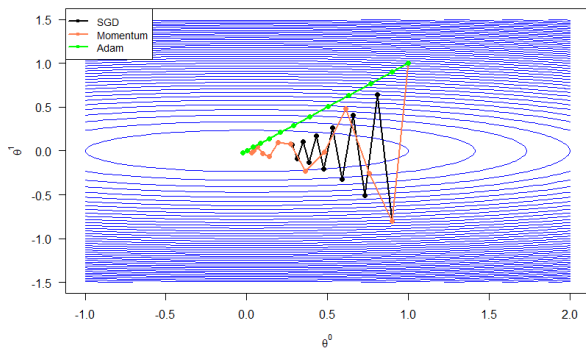


# ADAM

- Adaptive Moment Estimation also has adaptive step sizes
- Uses the 1st and 2nd moments of gradients
  - Keeps an exponentially decaying average of past gradients (1st moment)
  - Like RMSProp, stores an exponentially decaying average of past squared gradients (2nd moment)
  - Can be seen as combo of RMSProp + momentum.





# COMPARISON ON QUADRATIC FORM




SGD vs. SGD with Momentum vs. Adam on a quadratic form.

# ADAM

## Least-Squares:



●  $n = 5$  ●  $n = 3$  ●  $n = 2$  ●  $n = 1$



## Algorithm Adam

---

- 1: **require** Global step size  $\alpha$  (suggested default: 0.001)
  - 2: **require** Exponential decay rates for moment estimates,  $\rho_1$  and  $\rho_2$  in  $[0, 1)$  (suggested defaults: 0.9 and 0.999 respectively)
  - 3: **require** Small constant  $\beta$  (suggested default  $10^{-8}$ )
  - 4: **require** Initial parameters  $\theta$
  - 5: Initialize time step  $t = 0$
  - 6: Initialize 1st and 2nd moment variables  $\mathbf{s}^{[0]} = 0, \mathbf{r}^{[0]} = 0$
  - 7: **while** stopping criterion not met **do**
  - 8:    $t \leftarrow t + 1$
  - 9:   Sample a minibatch of  $m$  examples from the training set  $\{\tilde{\mathbf{x}}^{(1)}, \dots, \tilde{\mathbf{x}}^{(m)}\}$
  - 10:   Compute gradient estimate:  $\hat{\mathbf{g}}^{[t]} \leftarrow \frac{1}{m} \nabla_{\theta} \sum_i L(y^{(i)}, f(\tilde{\mathbf{x}}^{(i)} | \theta))$
  - 11:   Update biased first moment estimate:  $\mathbf{s}^{[t]} \leftarrow \rho_1 \mathbf{s}^{[t-1]} + (1 - \rho_1) \hat{\mathbf{g}}^{[t]}$
  - 12:   Update biased second moment estimate:  $\mathbf{r}^{[t]} \leftarrow \rho_2 \mathbf{r}^{[t-1]} + (1 - \rho_2) \hat{\mathbf{g}}^{[t]^2}$
  - 13:   Correct bias in first moment:  $\hat{\mathbf{s}} \leftarrow \frac{\mathbf{s}^{[t]}}{1 - \rho_1^t}$
  - 14:   Correct bias in second moment:  $\hat{\mathbf{r}} \leftarrow \frac{\mathbf{r}^{[t]}}{1 - \rho_2^t}$
  - 15:   Compute update for each entry  $i$ :  $\Delta \theta_i = -\alpha \frac{\hat{s}_i}{\sqrt{\hat{r}_i + \beta}}$
  - 16:   Apply update:  $\theta \leftarrow \theta + \Delta \theta$
  - 17: **end while**
- 



# ADAM / 3

- Initializes moment variables  $\mathbf{s}$  and  $\mathbf{r}$  with zero  $\Rightarrow$  Bias towards zero
- Indeed: Unrolling  $\mathbf{s}^{[t]}$  yields

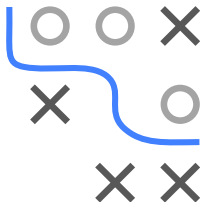
$$\mathbf{s}^{[0]} = 0$$

$$\mathbf{s}^{[1]} = \rho_1 \mathbf{s}^{[0]} + (1 - \rho_1) \hat{\mathbf{g}}^{[1]} = (1 - \rho_1) \hat{\mathbf{g}}^{[1]}$$

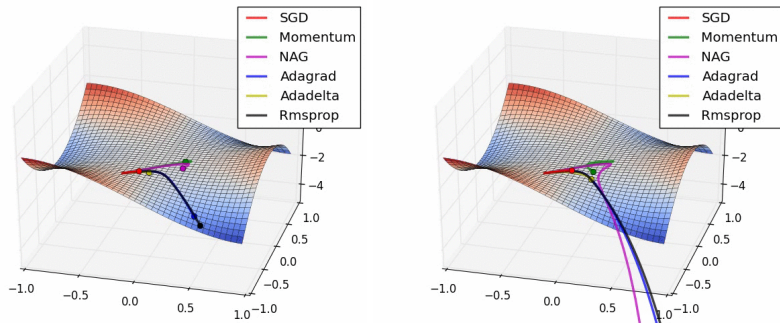
$$\mathbf{s}^{[2]} = \rho_1 \mathbf{s}^{[1]} + (1 - \rho_1) \hat{\mathbf{g}}^{[2]} = \rho_1 (1 - \rho_1) \hat{\mathbf{g}}^{[1]} + (1 - \rho_1) \hat{\mathbf{g}}^{[2]}$$

$$\mathbf{s}^{[3]} = \rho_1 \mathbf{s}^{[2]} + (1 - \rho_1) \hat{\mathbf{g}}^{[3]} = \rho_1^2 (1 - \rho_1) \hat{\mathbf{g}}^{[1]} + \rho_1 (1 - \rho_1) \hat{\mathbf{g}}^{[2]} + (1 - \rho_1) \hat{\mathbf{g}}^{[3]}$$

- Therefore:  $\mathbf{s}^{[t]} = (1 - \rho_1) \sum_{i=1}^t \rho_1^{t-i} \hat{\mathbf{g}}^{[i]}$ .
- Therefore:  $\mathbf{s}^{[t]}$  is a biased estimator of  $\hat{\mathbf{g}}^{[t]}$
- **Note:** Contributions of past  $\hat{\mathbf{g}}^{[i]}$  decreases rapidly and bias vanishes for  $t \rightarrow \infty$  ( $\rho_1^t \rightarrow 0$ )
- We correct for the bias by  $\hat{\mathbf{s}}^{[t]} = \frac{\mathbf{s}^{[t]}}{(1 - \rho_1^t)}$
- Analogously:  $\hat{\mathbf{r}}^{[t]} = \frac{\mathbf{r}^{[t]}}{(1 - \rho_2^t)}$



# COMPARISON OF OPTIMIZERS: ANIMATION



Credits: Dettmers (2015) and Radford

Comparison of SGD optimizers near saddle point.

**Left:** After start. **Right:** Later.

All methods accelerate compared to vanilla SGD.

Best is RMSProp, then AdaGrad. (Adam is missing here.)

